

15. Aggregation

Aggregation in MongoDB is a process of transforming and analyzing data stored in collections. It is similar to SQL's GROUP BY, SUM, JOIN, or WHERE clauses but is more flexible and powerful.

Uses of Aggregation:

1. Summarizing data (e.g., total sales, average age).
 2. Restructuring data (e.g., filtering, grouping, or reshaping).
 3. Analyzing data (e.g., performing statistical or conditional computations).
-

Aggregation Pipeline

The **aggregation pipeline** is a framework in MongoDB that allows the transformation of documents through a series of stages. Each stage performs a specific operation on the input documents and passes the transformed documents to the next stage.

Basic Syntax:

```
db.collection.aggregate([
  { stage1 },
  { stage2 },
  { stage3 }
]);
```

users Collection:

```
[
  {
    "_id": 1,
    "name": "Ajay",
    "age": 28,
    "gender": "Male",
    "hobbies": ["reading", "traveling", "yoga"],
    "scores": [40, 60, 80]
  },
  {
    "_id": 2,
    "name": "Sunita",
    "age": 25,
    "gender": "Female",
    "hobbies": ["painting", "gardening", "music"],
    "scores": [55, 45, 90]
  },
]
```

```
{
  "_id": 3,
  "name": "Ravi",
  "age": 35,
  "gender": "Male",
  "hobbies": ["gaming", "science fiction", "technology"],
  "scores": [70, 85, 95]
},
{
  "_id": 4,
  "name": "Madhavi",
  "age": 30,
  "gender": "Female",
  "hobbies": ["cooking", "dance", "fitness"],
  "scores": [30, 50, 70]
},
{
  "_id": 5,
  "name": "Sai",
  "age": 40,
  "gender": "Male",
  "hobbies": ["history", "painting", "walking"],
  "scores": [90, 95, 100]
}
]
```

orders Collection:

```
[
  {
    "_id": 101,
    "orderId": "ORD001",
    "customerId": 1,
    "amount": 500,
    "status": "completed",
    "orderDate": "2024-11-20T10:00:00Z",
    "items": [
      { "product": "Book", "quantity": 2, "price": 250 }
    ]
  },
  {
    "_id": 102,
    "orderId": "ORD002",
    "customerId": 2,
    "amount": 150,
    "status": "pending",
    "orderDate": "2024-11-21T12:00:00Z",
    "items": [
```

```

    { "product": "Pen", "quantity": 10, "price": 15 }
  ],
  {
    "_id": 103,
    "orderId": "ORD003",
    "customerId": 3,
    "amount": 300,
    "status": "completed",
    "orderDate": "2024-11-22T15:00:00Z",
    "items": [
      { "product": "Laptop", "quantity": 1, "price": 300 }
    ]
  },
  {
    "_id": 104,
    "orderId": "ORD004",
    "customerId": 1,
    "amount": 200,
    "status": "completed",
    "orderDate": "2024-11-23T14:00:00Z",
    "items": [
      { "product": "Headphones", "quantity": 1, "price": 200 }
    ]
  },
  {
    "_id": 105,
    "orderId": "ORD005",
    "customerId": 4,
    "amount": 100,
    "status": "canceled",
    "orderDate": "2024-11-24T16:00:00Z",
    "items": [
      { "product": "Notebook", "quantity": 5, "price": 20 }
    ]
  }
]

```

Key Aggregation Stages

1. \$match

Filters documents based on specified conditions. It's equivalent to the WHERE clause in SQL.

Example: Filter users with age greater than 25:

```

db.users.aggregate([
  { $match: { age: { $gt: 25 } } }
])

```

```
]);
```

2. \$group

Groups documents by a specified key and applies accumulator expressions to compute aggregated values for each group.

Example: Group users by gender and count the number of users in each group:

```
db.users.aggregate([
  { $group: { _id: "$gender", count: { $sum: 1 } } }
]);
```

Accumulators in \$group:

- \$sum: Calculates the sum of numeric values.
- \$avg: Computes the average of numeric values.
- \$min: Returns the minimum value.
- \$max: Returns the maximum value.
- \$push: Adds values to an array.
- \$addToSet: Adds unique values to an array.

3. \$project

Reshapes documents by including, excluding, or computing new fields.

Example: Include only name and age, and compute a new field isAdult:

```
db.users.aggregate([
  {
    $project: {
      _id: 0,
      name: 1,
      age: 1,
      isAdult: { $gte: ["$age", 18] }
    }
  }
]);
```

4. \$sort

Sorts documents by a specified field.

Example: Sort users by age in descending order:

```
db.users.aggregate([
  { $sort: { age: -1 } }
]);
```

5. \$limit

Limits the number of documents in the output.

Example: Retrieve the first 5 documents:

```
db.users.aggregate([
  { $limit: 5 }
]);
```

6. \$skip

Skips a specified number of documents.

Example: Skip the first 10 documents:

```
db.users.aggregate([
  { $skip: 10 }
]);
```

7. \$unwind

Deconstructs an array field into multiple documents, one for each element in the array.

Example: Unwind the hobbies array:

```
db.users.aggregate([
  { $unwind: "$hobbies" }
]);
```

Input:

```
{ "_id": 1, "name": "Alice", "hobbies": ["reading", "cycling"] }
```

Output:

```
{ "_id": 1, "name": "Alice", "hobbies": "reading" }
{ "_id": 1, "name": "Alice", "hobbies": "cycling" }
```

8. \$lookup

Performs a left outer join with another collection.

Example: Join orders with customers collection:

```
db.orders.aggregate([
```

```
{
  $lookup: {
    from: "customers",
    localField: "customerId",
    foreignField: "_id",
    as: "customerDetails"
  }
}
];
```

9. \$facet

Executes multiple pipelines in parallel and outputs the results in a single document.

Example: Calculate the total number of users and the top 5 users by age:

```
db.users.aggregate([
  {
    $facet: {
      totalUsers: [{ $count: "count" }],
      topUsers: [{ $sort: { age: -1 } }, { $limit: 5 }]
    }
  }
]);
```

10. \$bucket

Groups documents into ranges (buckets) defined by boundaries.

Example: Group users into age ranges (buckets):

```
db.users.aggregate([
  {
    $bucket: {
      groupBy: "$age",
      boundaries: [0, 18, 30, 50],
      default: "Other",
      output: {
        count: { $sum: 1 },
        names: { $push: "$name" }
      }
    }
  }
]);
```

11. \$filter

Filters elements of an array based on a condition.

Example: Filter scores array to retain only values greater than 50:

```
db.users.aggregate([
  {
    $project: {
      name: 1,
      passingScores: {
        $filter: {
          input: "$scores",
          as: "score",
          cond: { $gt: ["$$score", 50] }
        }
      }
    }
  }
]);
```

Pipeline Expressions

MongoDB's aggregation framework supports expressions to compute values. Some common expressions include:

- **\$add:** Adds numbers.
- **\$subtract:** Subtracts numbers.
- **\$multiply:** Multiplies numbers.
- **\$divide:** Divides numbers.
- **\$concat:** Concatenates strings.
- **\$arrayElemAt:** Returns an element from an array.

Performance Tips

1. **Use \$match Early:**
 - Filter as many documents as possible early in the pipeline to reduce processing.
2. **Optimize with Indexes:**
 - Fields used in \$match and \$sort stages should be indexed for faster performance.
3. **Avoid \$unwind on Large Arrays:**
 - \$unwind can be resource-intensive if applied to arrays with many elements.
4. **Use \$facet Sparingly:**
 - It runs parallel pipelines but can consume more memory.
5. **Analyze with explain():**
 - Use the explain() method to analyze query performance.

interview questions and answers

1. What is the Aggregation Framework in MongoDB?

Answer: The aggregation framework is a set of tools in MongoDB used to process and transform data stored in collections. It allows developers to:

- Filter, group, and sort data.
- Compute aggregated values like sums, averages, and counts.
- Perform operations like joining collections and filtering array elements.

It is similar to SQL's GROUP BY, SUM, and WHERE clauses.

2. What is an aggregation pipeline?

Answer: An aggregation pipeline is a sequence of stages, where each stage processes documents and passes the output to the next stage. Each stage performs a specific operation such as filtering, grouping, or reshaping documents.

Example Syntax:

```
db.collection.aggregate([
  { $match: { age: { $gt: 25 } } },
  { $group: { _id: "$gender", avgAge: { $avg: "$age" } } }
]);
```

3. What are the main stages in an aggregation pipeline?

Answer: Key aggregation stages include:

- **\$match:** Filters documents based on conditions.
- **\$group:** Groups documents by a field and performs aggregations.
- **\$project:** Reshapes documents by including, excluding, or adding fields.
- **\$sort:** Sorts documents by one or more fields.
- **\$limit:** Limits the number of output documents.
- **\$skip:** Skips a specified number of documents.
- **\$unwind:** Deconstructs arrays into multiple documents.
- **\$lookup:** Joins two collections.

4. What is the \$group stage used for?

Answer: The \$group stage groups documents by a specified field and applies accumulator expressions to compute values for each group.

Example: Calculate the total number of users in each gender group:

```
db.users.aggregate([
  { $group: { _id: "$gender", userCount: { $sum: 1 } } }
]);
```

5. What is the difference between \$match and \$filter?

Answer:

- **\$match:**
 - Filters documents at the pipeline level.
 - Used to filter documents based on conditions (like a WHERE clause in SQL).
 - Operates on entire documents.

```
{ $match: { age: { $gt: 30 } } }
```

\$filter:

- Filters elements within an array.
- Used within stages like \$project or \$addFields

```
{
  $project: {
    passingScores: {
      $filter: {
        input: "$scores",
        as: "score",
        cond: { $gt: ["$$score", 50] }
      }
    }
  }
}
```

6. What is the \$lookup stage, and how is it used?

Answer: The \$lookup stage is used to perform a left outer join between two collections. It allows you to merge fields from another collection into the output.

Example: Join the orders collection with the users collection:

```
db.orders.aggregate([
  {
    $lookup: {
      from: "users",
      localField: "customerId",
      foreignField: "_id",
      as: "userDetails"
    }
  }
]);
```

7. Explain the \$unwind stage with an example.

Answer: The \$unwind stage deconstructs an array field from a document into multiple documents, one for each element in the array.

Example: Input:

```
{ "_id": 1, "name": "Ajay", "hobbies": ["reading", "yoga"] }
```

Query:

```
db.users.aggregate([
  { $unwind: "$hobbies" }
]);
```

Output:

```
{ "_id": 1, "name": "Ajay", "hobbies": "reading" }
{ "_id": 1, "name": "Ajay", "hobbies": "yoga" }
```

8. What is \$facet, and how is it used?

Answer: The \$facet stage allows you to run multiple aggregation pipelines in parallel and output their results as a single document.

Example: Find the total number of users and the top 5 oldest users:

```
db.users.aggregate([
  {
    $facet: {
      totalUsers: [{ $count: "count" }],
      topUsers: [{ $sort: { age: -1 } }, { $limit: 5 }]
    }
  }
]);
```

9. What is \$bucket, and how is it different from \$group?

Answer:

- **\$bucket:**
 - Groups documents into defined ranges (buckets) based on a field or expression.
 - Requires specifying boundaries.

Example: Group users into age ranges:

```
db.users.aggregate([
  {
```

```
$bucket: {  
  groupBy: "$age",  
  boundaries: [20, 30, 40, 50],  
  output: { count: { $sum: 1 } }  
}  
}  
]);
```

- **\$group:**
 - Groups documents by a specific field or computed expression.
 - Does not require fixed ranges.
-

10. How does \$sort work in aggregation?

Answer: The \$sort stage orders documents based on a field in ascending (1) or descending (-1) order.

Example: Sort users by age in descending order:

```
db.users.aggregate([  
  { $sort: { age: -1 } }  
]);
```